

Benchmark Performance Evaluation

A20375099 Yuan-An Liu, A20382007 Sharel Clavy Pereira

1. Introduction

All the experiments have been carried out on on Chameleon KVM on 8 virtual core instances with 16GB of RAM and 160GB disk. GPU benchmarks have been done on a Tesla K80 GPU on Chameleon baremetal instance. All the tables and graphs are properly marked with units and brief description for easier understanding of what is going on in the experiments.

```
[cc@pal-liu-pereira CPU]$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
CPU(s):                8
On-line CPU(s) list:   0-7
Thread(s) per core:    1
Core(s) per socket:    1
Socket(s):             8
NUMA node(s):          1
Vendor ID:             GenuineIntel
CPU family:            6
Model:                 60
Model name:            Intel Core Processor (Haswell)
Stepping:              1
CPU MHz:               2299.998
BogoMIPS:              4599.99
Hypervisor vendor:     KVM
Virtualization type:   full
L1d cache:             32K
L1i cache:             32K
L2 cache:              4096K
NUMA node0 CPU(s):    0-7
Flags:                 fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov
v pat pse36 clflush mmx fxsr sse sse2 ss syscall nx pdpe1gb rdtscp lm constant_tsc
rep_good nopl eagerfpu pni pclmulqdq ssse3 fma cx16 pcid sse4_1 sse4_2 x2apic movbe
popcnt tsc_deadline_timer aes xsave avx f16c rdrand hypervisor lahf_lm abm fsgsbase
e bmi1 avx2 smep bmi2 erms invpcid xsaveopt
```

2. Experimental Results


In this section we discuss the results gathered for different benchmarks along with the theoretical results and the trade offs achieved.


2.1 CPU Benchmark

We are measuring double precision floating point operations and integer operations, per second, in terms of GFLOPS and GIOPS using 1,2,4,8 threads.

Average and **Standard deviation** for are as follows:


| | <i>No. of Threads</i> | <i>Average</i> | <i>Standard Deviation</i> |
|---------------|-----------------------|----------------|---------------------------|
| GFLOPS | 1 | 0.003867 | 0.00019798 |
| | 2 | 0.00363 | 0.000140105 |
| | 4 | 0.003546 | 0.000141643 |
| | 8 | 0.003404 | 4.48876E-05 |
| GIOPS | 1 | 0.004759 | 9.88786E-05 |
| | 2 | 0.004709 | 4.02315E-05 |
| | 4 | 0.004387 | 4.48234E-05 |
| | 8 | 0.004173 | 0.000377654 |






Results of 600 samples for FLOPS/IOPS with 8 threads for a 10min period:

We can notice that results are almost similar for quite a lot of time throughout the time period of 600s except for a few dips which might be caused because of different actions taking place in the operating system during the process. No process was invoked by any external process during this experiment.





We have carried out GFLOPS and GIOPS calculation for multiple thread processes. Above plots and table define the results obtained. We notice that an average of 0.003404GFLOPS and 0.004173GIOPS is achieved with multiple operations.

Linpack Benchmark results are as below:

We also notice that in the Linpack results we have 1.528e-01GFLOPS . We also can conclude that results varies with multiple operations and multithreading in place. As we scale up, relative value of operations completed varies. Full Result please visit file at CPU/linpack.txt

| T/V | N | NB | P | Q | Time | Gflops |
|----------|----|----|---|---|------|-----------|
| WR00R2R2 | 35 | 4 | 4 | 1 | 0.00 | 1.179e-01 |

HPL_pdgesv() start time Tue Sep 26 19:53:52 2017

HPL_pdgesv() end time Tue Sep 26 19:53:52 2017

--VVV--VVV--VVV--VVV--VVV--VVV--VVV--VVV--VVV--VVV--VVV--VVV--VVV--VVV--VVV--

Max aggregated wall time rfact . . . : 0.00
+ Max aggregated wall time pfact . . . : 0.00
+ Max aggregated wall time mxswp . . . : 0.00
Max aggregated wall time update . . . : 0.00
+ Max aggregated wall time laswp . . . : 0.00
Max aggregated wall time up tr sv . . : 0.00

||Ax-b||_oo/(eps*(||A||_oo*||x||_oo+||b||_oo)*N)= 0.0171791 PASSED

Column=000000004 Fraction=11.4% Gflops=3.792e-01
Column=000000008 Fraction=22.9% Gflops=3.361e-01
Column=000000012 Fraction=34.3% Gflops=3.102e-01
Column=000000016 Fraction=45.7% Gflops=2.475e-01
Column=000000020 Fraction=57.1% Gflops=2.251e-01
Column=000000024 Fraction=68.6% Gflops=2.007e-01
Column=000000028 Fraction=80.0% Gflops=1.783e-01
Column=000000032 Fraction=91.4% Gflops=1.528e-01

| T/V | N | NB | P | Q | Time | Gflops |
|----------|----|----|---|---|------|-----------|
| WR00R2R4 | 35 | 4 | 4 | 1 | 0.00 | 1.399e-01 |

HPL_pdgesv() start time Tue Sep 26 19:53:52 2017

HPL_pdgesv() end time Tue Sep 26 19:53:52 2017

--VVV--VVV--VVV--VVV--VVV--VVV--VVV--VVV--VVV--VVV--VVV--VVV--VVV--VVV--VVV--

Max aggregated wall time rfact . . . : 0.00
+ Max aggregated wall time pfact . . . : 0.00
+ Max aggregated wall time mxswp . . . : 0.00
Max aggregated wall time update . . . : 0.00
+ Max aggregated wall time laswp . . . : 0.00
Max aggregated wall time up tr sv . . : 0.00

||Ax-b||_oo/(eps*(||A||_oo*||x||_oo+||b||_oo)*N)= 0.0168365 PASSED

Finished 864 tests with the following results:
864 tests completed and passed residual checks,
0 tests completed and failed residual checks,
0 tests skipped because of illegal input values.


End of Tests.

2.2 GPU Benchmark

Results of multiple experiments are as follows: (Mapping one per thread per core of a 2496 core K80 GPU). Experiment is measured for double precision floating point operations , integer operations, half-precision operations and quarter-precision operations.

| | | | |
|--|-------------------------|-----------------------|---|
| Site: tacc | Cluster: chameleon | Platform Type: x86_64 | # CPUs: 2 |
| # of Threads: 48 | RAM Size: 128 GiB | Node Type: gpu_k80 | Wattmeter: No |
| Version: 86f6934b0783bd209daace2bf4805b8d1614f689 | | | |
| Bios | | | |
| Release Date: 03/09/2015 | Vendor: 1.2 | Version: Dell Inc. | |
| Chassis | | | |
| Manufacturer: Dell Inc. | Name: PowerEdge R730 | Serial: 1M2KD42 | |
| GPU | | | |
| GPU: No | | | |
| Network Adapters | | | More ▶ |
| Operating System | | | |
| Kernel: | Name: | Version: | |
| Processor | | | |
| Cache L1d: 32768 | Cache L1i: 32768 | Cache L2: 262144 | Cache L3: 31457280 |
| Clock Speed: 3100000000 | Instruction Set: x86-64 | Model: Intel Xeon | Other Description: Intel(R) Xeon(R) CPU E5-2670 v3 @ 2.30GHz |
| Vendor: Intel | Version: E5-2670 v3 | | |
| Storage Devices | | | More ▶ |
| Supported Job Types | | | |
| Best Effort: No | Deploy: No | Virtual: ivt | |

| Experiment No. | GFLOPS | GIOPS | GHOPS | GQOPS |
|----------------|---------|---------|---------|---------|
| 2496_Exp1 | 0.00138 | 0.00151 | 0.00147 | 0.00128 |
| 2496_Exp2 | 0.00151 | 0.0015 | 0.00127 | 0.00147 |
| 2496_Exp3 | 0.00143 | 0.0014 | 0.00151 | 0.00153 |



For a **NVIDIA Tesla K80** processor with dual processing units the theoretical speed for a double precision floating point performance is 1.87 TFLOPS. We have achieved about an average of 0.0014 GFLOPS in our experiment. Since we carried out the experiments by assigning only one thread per core(for 2496 cores) of the GPU , the figurative performance obtained is relative to that of the theoretical peak value. Hence we can conclude that the experimental value obtained is debatable with that of theoretical value.

Linpack Benchmark results are as below:

```
-----
- The matrix A is randomly generated for each test.
- The following scaled residual check will be computed:
  ||Ax-b||_oo / ( eps * ( ||x||_oo * ||A||_oo + ||b||_oo ) * N )
- The relative machine precision (eps) is taken to be      1.110223e-16
- Computational tests pass if scaled residuals are less than 16.0

Column=000000001 Fraction= 3.4% Gflops=8.200e-04|
```

We notice that Linpack has an average of 8.200e-04 GFLOPS for the operations on the GPU. We have derived slightly different figure since we are running only one thread per core of the K80 gpu. However if we change the metrics and add in multiple threads per core with spawning multiple blocks we can probably get a different figure for comparisons.

2.3 Memory Benchmark

A program is created to measure the speed of the memory by performing various operations like read+write, sequential read access and random read access. Optimum concurrency is achieved by passing multiple threads and blocks of varying size (1B, 1KB, 1MB, 10MB), on a 20MB file.


```
[cc@pal-liu-pereira ~]$ sudo dmidecode --type memory
# dmidecode 3.0
Getting SMBIOS data from sysfs.
SMBIOS 2.4 present.


Handle 0x1000, DMI type 16, 15 bytes
Physical Memory Array
    Location: Other
    Use: System Memory
    Error Correction Type: Multi-bit ECC
    Maximum Capacity: 16 GB
    Error Information Handle: Not Provided
    Number Of Devices: 1

Handle 0x1100, DMI type 17, 21 bytes
Memory Device
    Array Handle: 0x1000
    Error Information Handle: 0x0000
    Total Width: 64 bits
    Data Width: 64 bits
    Size: 16384 MB
    Form Factor: DIMM
    Set: None
    Locator: DIMM 0
    Bank Locator: Not Specified
    Type: RAM
    Type Detail: None
```



Throughput is measured in **MB/sec** and Latency in **ms** for Memory benchmarks.

| Read + Write | No. of Threads | 1B | 1KB | 1MB | 10MB |
|--------------|----------------|----------|----------|----------|----------|
| Throughput | 1 | 9.06E+06 | 5.42E+08 | 9.21E+08 | 1.55E+09 |
| | 2 | 5.33E+06 | 8.95E+07 | 8.40E+08 | 1.51E+09 |
| | 4 | 1.73E+06 | 3.87E+07 | 6.55E+08 | 1.37E+09 |
| | 8 | 650956 | 2.16E+07 | 3.54E+08 | 6.53E+08 |
| Latency | 1 | 1.10E-10 | 1.84E-09 | 1.09E-06 | 6.46E-06 |
| | 2 | 1.88E-10 | 1.12E-08 | 1.19E-06 | 6.62E-06 |
| | 4 | 5.77E-10 | 2.58E-08 | 1.53E-06 | 7.32E-06 |
| | 8 | 1.54E-09 | 4.64E-08 | 2.83E-06 | 1.53E-05 |







| Sequential Read | No. of Threads | 1B | 1KB | 1MB | 10MB |
|-----------------|----------------|----------|----------|----------|----------|
| Throughput | 1 | 1.03E+08 | 4.39E+10 | 5.36E+10 | 9.00E+10 |
| | 2 | 9.41E+07 | 2.99E+10 | 3.91E+10 | 5.53E+10 |
| | 4 | 9.18E+07 | 1.70E+10 | 3.32E+10 | 3.05E+10 |
| | 8 | 8.94E+07 | 9.79E+09 | 1.54E+10 | 1.16E+10 |
| Latency | 1 | 9.67E-12 | 2.28E-11 | 1.87E-08 | 1.11E-07 |
| | 2 | 1.06E-11 | 3.34E-11 | 2.55E-08 | 1.81E-07 |
| | 4 | 1.09E-11 | 5.87E-11 | 3.01E-08 | 3.27E-07 |
| | 8 | 1.12E-11 | 1.02E-10 | 6.50E-08 | 8.62E-07 |




| Random Read | No. of Threads | 1B | 1KB | 1MB | 10MB |
|-------------|----------------|----------|----------|----------|----------|
| Throughput | 1 | 9.29E+07 | 2.26E+10 | 6.20E+10 | 5.75E+10 |
| | 2 | 9.80E+07 | 2.20E+10 | 3.42E+10 | 4.74E+10 |
| | 4 | 9.36E+07 | 1.72E+10 | 2.49E+10 | 2.88E+10 |
| | 8 | 6.94E+07 | 5.80E+09 | 1.42E+10 | 6.93E+09 |


| | | | | | |
|---------|---|----------|----------|----------|----------|
| Latency | 1 | 1.08E-11 | 4.42E-11 | 1.61E-08 | 1.74E-07 |
| | 2 | 1.02E-11 | 4.55E-11 | 2.93E-08 | 2.11E-07 |
| | 4 | 1.07E-11 | 5.82E-11 | 4.02E-08 | 3.47E-07 |
| | 8 | 1.44E-11 | 1.72E-10 | 7.02E-08 | 1.44E-06 |





Average and Standard deviation for are as follows:

| Read + Write (Avg) | No. of Threads | 1B | 1KB | 1MB | 10MB |
|--------------------|----------------|----------|----------|----------|----------|
| Throughput | 1 | 9.33E+06 | 5.49E+08 | 9.12E+08 | 1.55E+09 |
| | 2 | 6.54E+06 | 9.43E+07 | 8.51E+08 | 1.48E+09 |
| | 4 | 1.62E+06 | 3.63E+07 | 6.73E+08 | 1.28E+09 |
| | 8 | 6.82E+05 | 2.28E+07 | 3.66E+08 | 6.56E+08 |
| Latency | 1 | 1.07E-10 | 1.82E-09 | 1.10E-06 | 6.45E-06 |
| | 2 | 1.56E-10 | 1.06E-08 | 1.17E-06 | 6.79E-06 |
| | 4 | 6.19E-10 | 2.76E-08 | 1.49E-06 | 7.81E-06 |
| | 8 | 1.47E-09 | 4.41E-08 | 2.74E-06 | 1.52E-05 |





| Read + Write (SD) | No. of Threads | 1B | 1KB | 1MB | 10MB |
|-------------------|----------------|----------|----------|----------|----------|
| Throughput | 1 | 2.54E+05 | 6.72E+06 | 2.91E+07 | 7.04E+07 |
| | 2 | 1.11E+06 | 5.54E+06 | 1.11E+07 | 1.33E+08 |
| | 4 | 9.85E+04 | 2.11E+06 | 4.35E+07 | 7.70E+07 |
| | 8 | 4.00E+04 | 2.00E+06 | 1.37E+07 | 9.51E+06 |
| Latency | 1 | 2.94E-12 | 2.23E-11 | 3.55E-08 | 2.92E-07 |
| | 2 | 2.82E-11 | 6.14E-10 | 1.54E-08 | 6.29E-07 |
| | 4 | 3.64E-11 | 1.55E-09 | 9.34E-08 | 4.62E-07 |
| | 8 | 8.40E-11 | 3.69E-09 | 1.01E-07 | 2.19E-07 |



From the experiment we have achieved an average of $6.82E+05$ MB/s for a 1B Read + Write operations running on 8 threads while the average latency is $1.47E-09$ ms. We have performed experiments by scaling up multiple threads and size of data in memory across multiple operations. Results and graphs are displayed as above for each operation.

STREAM Benchmark results are as below:

According to Linpack results we can achieve an average throughput of 13983.5MB/s. However throughput varies with different operations and with different number of threads. When multiple threads try to write the same file, they have to wait for previous one to complete and hence the throughput differs in these cases.

The result is stored in Memory/stream.txt

STREAM version \$Revision: 5.10 \$

This system uses 8 bytes per array element.

Array size = 10000000 (elements), Offset = 0 (elements)

Memory per array = 76.3 MiB (= 0.1 GiB).

Total memory required = 228.9 MiB (= 0.2 GiB).

Each kernel will be executed 10 times.

The *best* time for each kernel (excluding the first iteration)
will be used to compute the reported bandwidth.

Your clock granularity/precision appears to be 1 microseconds.

Each test below will take on the order of 11909 microseconds.

(= 11909 clock ticks)

Increase the size of the arrays if this shows that
you are not getting at least 20 clock ticks per test.

WARNING -- The above is only a rough guideline.

For best results, please be sure you know the
precision of your system timer.

| Function | Best Rate MB/s | Avg time | Min time | Max time |
|----------|----------------|----------|----------|----------|
| Copy: | 12533.4 | 0.013056 | 0.012766 | 0.013201 |
| Scale: | 12477.5 | 0.013052 | 0.012823 | 0.013392 |
| Add: | 13983.5 | 0.017517 | 0.017163 | 0.017954 |
| Triad: | 13970.5 | 0.017706 | 0.017179 | 0.018359 |

Solution Validates: avg error less than 1.000000e-13 on all three arrays

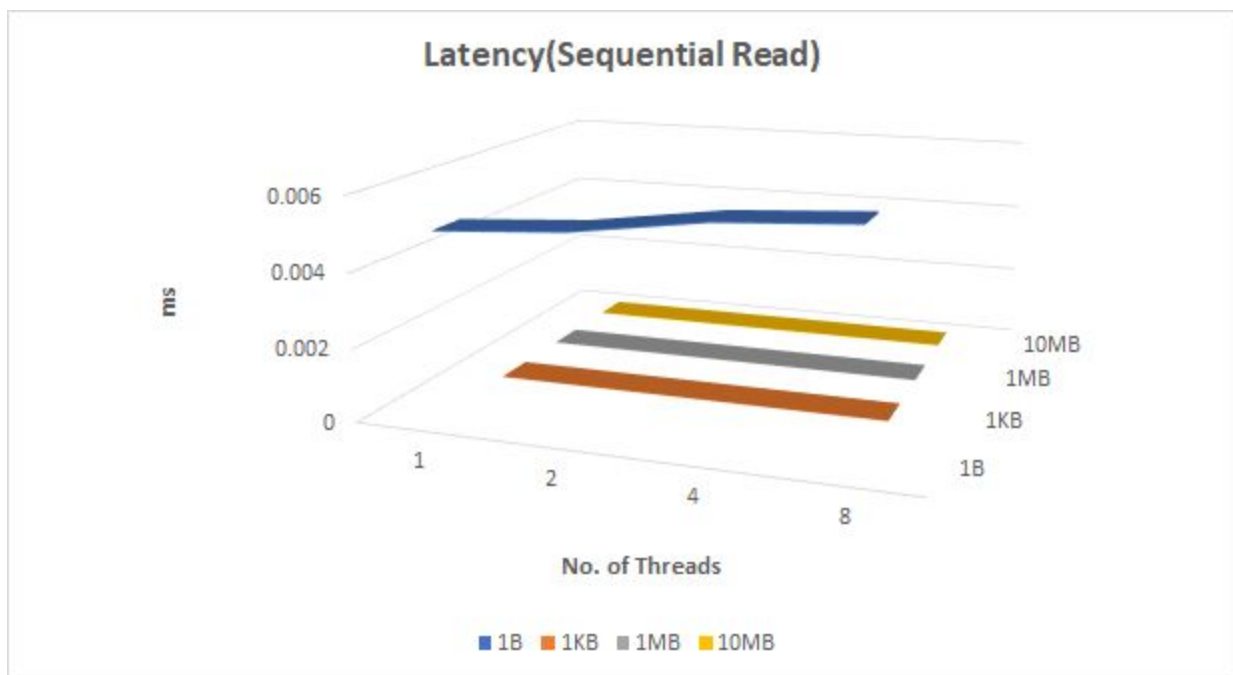
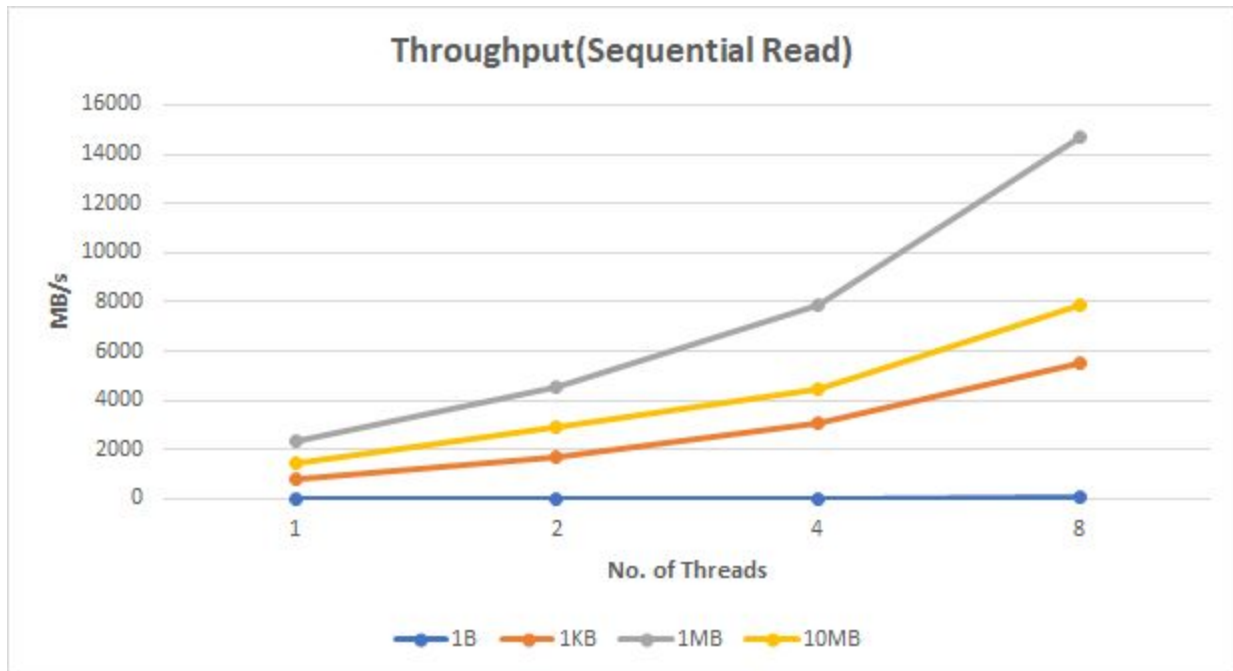
2.4 Disk Benchmark

In this section we are measuring the speed of the disk using different parameter space like read+write, sequential read access and random read access. Concurrency is achieved by passing multiple threads with blocks of varying size (1B, 1KB, 1MB, 10MB), using a 20MB file.

```
[cc@pal-liu-pereira ~]$ lsblk --output NAME,FSTYPE,UUID,RO,RM,MODEL,SIZE,MODE,ROTA,RQ-SIZE,TYPE,DISC-MAX
NAME FSTYPE UUID RO RM MODEL SIZE MODE ROTA RQ-SIZE TYPE DISC-MAX
vda
└─vda1 xfs d613911c-f1e0-4243-903e-535ec22d4891 0 0 160G brw-rw---- 1 128 disk 0B
[cc@pal-liu-pereira ~]$
```


Throughput is measured in MB/sec and Latency in ms for Disk benchmarks.

| Sequential Read | No. of Threads | 1B | 1KB | 1MB | 10MB |
|-----------------|----------------|-------------|-------------|-------------|-------------|
| Throughput | 1 | 3.984067734 | 811.6934212 | 2333.734316 | 1450.112018 |
| | 2 | 7.612959007 | 1656.305569 | 4573.442373 | 2886.601401 |
| | 4 | 13.87236956 | 3090.466594 | 7864.811551 | 4474.282209 |
| | 8 | 26.73232203 | 5514.558154 | 14658.9917 | 7831.677812 |
| Latency | 1 | 0.005019995 | 2.41E-08 | 8.17E-12 | 1.32E-12 |
| | 2 | 0.005254199 | 2.36E-08 | 8.34E-12 | 1.32E-12 |
| | 4 | 0.005766859 | 2.53E-08 | 9.70E-12 | 1.71E-12 |
| | 8 | 0.005985264 | 2.83E-08 | 1.04E-11 | 1.95E-12 |





| Random Read | No. of Threads | 1B | 1KB | 1MB | 10MB |
|-------------|----------------|-------------|-------------|-------------|-------------|
| Throughput | 1 | 0.429725497 | 201.205225 | 2628.092359 | 1461.862921 |
| | 2 | 0.767411181 | 422.645677 | 4833.81814 | 2804.054018 |
| | 4 | 1.441065593 | 807.5267435 | 8545.198767 | 5161.346849 |
| | 8 | 2.565619035 | 1557.465774 | 14772.57727 | 9473.568424 |
| Latency | 1 | 0.046541339 | 9.71E-08 | 7.26E-12 | 1.30E-12 |
| | 2 | 0.052123296 | 9.24E-08 | 7.89E-12 | 1.36E-12 |
| | 4 | 0.055514475 | 9.67E-08 | 8.93E-12 | 1.48E-12 |
| | 8 | 0.062363117 | 1.00E-07 | 1.03E-11 | 1.61E-12 |







| Read and Write | No. of Threads | 1B | 1KB | 1MB | 10MB |
|----------------|----------------|-------------|-------------|-------------|-------------|
| Throughput | 1 | 1.816034767 | 268.8458223 | 12.73613664 | 68.99337669 |
| | 2 | 3.622946185 | 115.8537101 | 61.14805387 | 400.5408916 |
| | 4 | 4.697484638 | 188.803873 | 26.99696612 | 210.244754 |
| | 8 | 13.34959315 | 69.91418464 | 175.7621767 | 602.5346569 |
| Latency | 1 | 0.011013005 | 7.26E-08 | 1.50E-09 | 2.76E-11 |
| | 2 | 0.011040738 | 3.37E-07 | 6.24E-10 | 9.52E-12 |
| | 4 | 0.017030391 | 4.14E-07 | 2.83E-09 | 3.63E-11 |
| | 8 | 0.011985384 | 2.23E-06 | 8.68E-10 | 2.53E-11 |



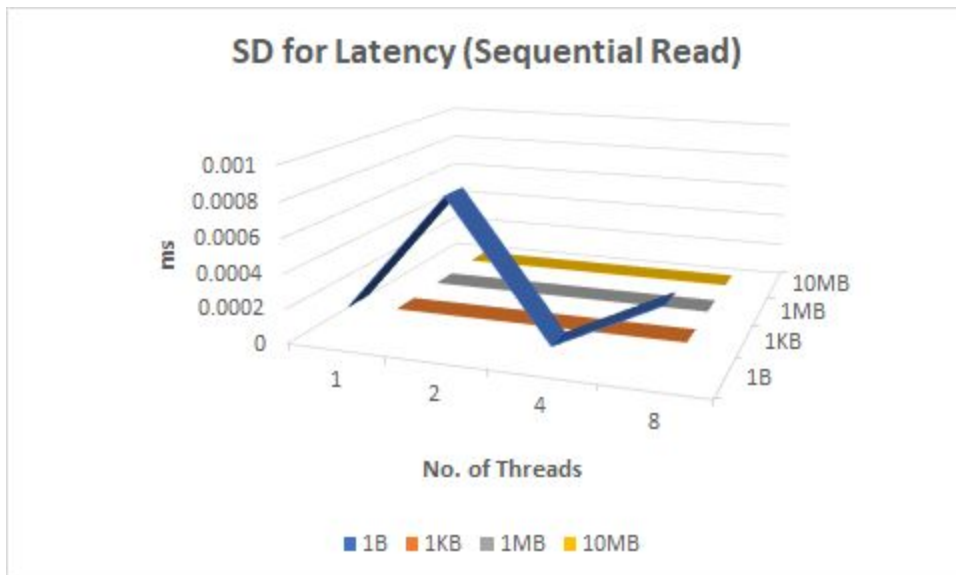
Average and Standard deviation for are as follows:

| Sequential Read (Avg) | NO. of Threads | 1B | 1KB | 1MB | 10MB |
|-----------------------|----------------|-------------|-------------|----------|-----------|
| Throughput | 1 | 3.823210718 | 829.4546963 | 2308.53 | 1524.04 |
| | 2 | 6.903690689 | 1614.989205 | 4376.038 | 2819.4395 |
| | 4 | 13.91614839 | 3104.835034 | 8157.283 | 4861.1999 |
| | 8 | 26.10821944 | 4912.340177 | 13701.9 | 8112.3075 |
| Latency | 1 | 0.005236138 | 2.36E-08 | 8.27E-12 | 1.25E-12 |
| | 2 | 0.005872229 | 2.42E-08 | 8.73E-12 | 1.35E-12 |
| | 4 | 0.005749992 | 2.52E-08 | 9.36E-12 | 1.58E-12 |
| | 8 | 0.006144557 | 3.32E-08 | 1.12E-11 | 1.89E-12 |





| Sequential Read (SD) | NO. of Threads | 1B | 1KB | 1MB | 10MB |
|----------------------|----------------|-------------|-------------|----------|-----------|
| Throughput | 1 | 0.144756525 | 17.27838132 | 83.94254 | 66.610892 |
| | 2 | 0.941942917 | 37.44901186 | 186.5776 | 101.1686 |
| | 4 | 0.254330658 | 24.46893778 | 253.6396 | 336.41064 |
| | 8 | 1.617278947 | 1150.050556 | 1464.028 | 653.17958 |
| Latency | 1 | 0.000195435 | 4.92E-10 | 3.05E-13 | 5.57E-14 |
| | 2 | 0.000860673 | 5.57E-10 | 3.69E-13 | 4.96E-14 |
| | 4 | 0.000104622 | 1.97E-10 | 2.96E-13 | 1.13E-13 |
| | 8 | 0.000392808 | 8.98E-09 | 1.28E-12 | 1.46E-13 |



We have achieved an average throughput of 13701.9MB/s for a 1MB Sequential Read operation with an average latency of 1.28E-12ms with multithreading of 8 threads. From the results we also can deduce that higher throughput is achieved for higher number of threads while sequentially reading a block of data. We also notice that the latency reduces as multiple threads are spawned for a particular operation.

Though it seems to be performing faster, average SSD speed might be faster than this for the given configurations. We can thus say that HDD is being used here.

iozone Benchmark results are as below:

Below results shows the output for the iozone benchmark. We see that a average throughput of 1306.3MB/s is achieved in this case. Also the throughput decreases if multiple processes are trying to write to the same file because latency increases in this case. Threads will have to wait for other threads to finish the process to start writing or any other data change operation to the file.


Full Result please visit file at Disk/iozone_output.xls


| | | | | | | | | | | | | | |
|---|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|---------|
| ./iozone -a -b output.xls | | | | | | | | | | | | | |
| The top row is records sizes, the left column is file sizes | | | | | | | | | | | | | |
| Writer Report | | | | | | | | | | | | | |
| | 4 | 8 | 16 | 32 | 64 | 128 | 256 | 512 | 1024 | 2048 | 4096 | 8192 | 16384 |
| 64 | 703068 | 877797 | 889431 | 969761 | 955947 | | | | | | | | |
| 128 | 800337 | 978253 | 1152564 | 955616 | 1155043 | 1102844 | | | | | | | |
| 256 | 937924 | 1219045 | 1108314 | 1248819 | 1272498 | 1260547 | 1230218 | | | | | | |
| 512 | 977103 | 1190529 | 1237168 | 1147892 | 1218910 | 1344020 | 1234323 | 1340664 | | | | | |
| 1024 | 1017811 | 1194770 | 1306332 | 1416636 | 1439908 | 1354535 | 1254808 | 1317957 | 1438461 | | | | |
| 2048 | 1060642 | 1272054 | 1376193 | 1411470 | 1481072 | 1473198 | 1334295 | 2306143 | 2334979 | 3003677 | | | |
| 4096 | 1509149 | 2183288 | 2858164 | 2957047 | 2950952 | 3105653 | 2981164 | 2856263 | 1476592 | 1463386 | 1407516 | | |
| 8192 | 937949 | 1411726 | 1681761 | 1850102 | 1884087 | 1405604 | 1486996 | 1684234 | 1986019 | 1365551 | 1612617 | 1229689 | |
| 16384 | 1068270 | 1095256 | 1159767 | 1545375 | 1614808 | 1500366 | 1485543 | 1592837 | 1595166 | 1490117 | 1757000 | 1457903 | 1338131 |
| 32768 | 0 | 0 | 0 | 0 | 1387658 | 1472600 | 1502363 | 2315782 | 1473310 | 1471858 | 1472505 | 1417118 | 1404720 |
| 65536 | 0 | 0 | 0 | 0 | 1520773 | 1582880 | 1474707 | 1567731 | 1633331 | 2108496 | 1667373 | 1698929 | 2004216 |
| 131072 | 0 | 0 | 0 | 0 | 1586808 | 2055547 | 1821023 | 1885878 | 1934011 | 1858150 | 2019974 | 1874094 | 1793423 |
| 262144 | 0 | 0 | 0 | 0 | 2063606 | 2229587 | 2114897 | 1984152 | 1637970 | 2234798 | 1833776 | 1852318 | 1529570 |
| 524288 | 0 | 0 | 0 | 0 | 1852856 | 2192234 | 2163322 | 2140021 | 2146197 | 2178469 | 2171199 | 2155417 | 1972402 |
| Re-writer Report | | | | | | | | | | | | | |
| | 4 | 8 | 16 | 32 | 64 | 128 | 256 | 512 | 1024 | 2048 | 4096 | 8192 | 16384 |
| 64 | 1933893 | 1599680 | 1828508 | 2067979 | 2133730 | | | | | | | | |
| 128 | 1294270 | 1858934 | 2326073 | 2286447 | 2558895 | 2608629 | | | | | | | |
| 256 | 1826695 | 2345409 | 2533571 | 2692393 | 2607399 | 2539563 | 2754556 | | | | | | |
| 512 | 1909204 | 1970519 | 2572435 | 2449212 | 2860238 | 2926501 | 2796910 | 2845081 | | | | | |
| 1024 | 1822368 | 2311849 | 2348509 | 2716948 | 2553783 | 2604895 | 2744728 | 2625597 | 2625597 | | | | |
| 2048 | 1859773 | 2255875 | 2656292 | 2798201 | 2779189 | 2921939 | 2727131 | 4751580 | 5238385 | 6132138 | | | |
| 4096 | 3767675 | 5051705 | 3761077 | 5704175 | 6043279 | 6195843 | 5978090 | 1970212 | 2892815 | 2957556 | 2919360 | | |
| 8192 | 2084454 | 3030813 | 3123099 | 3953573 | 1205232 | 2842277 | 2966955 | 3402019 | 4037664 | 3161607 | 2925775 | 2080793 | |
| 16384 | 1822903 | 1680930 | 1872016 | 2677993 | 3034132 | 2819514 | 3196878 | 2929368 | 2728933 | 2928868 | 2099470 | 2629524 | 2278436 |
| 32768 | 0 | 0 | 0 | 0 | 2309905 | 2437906 | 2840011 | 2146752 | 2385396 | 3229313 | 2250706 | 2222751 | 2241456 |
| 65536 | 0 | 0 | 0 | 0 | 2291626 | 2335812 | 2321352 | 2359876 | 2432659 | 2398652 | 2342340 | 3422037 | 2112839 |
| 131072 | 0 | 0 | 0 | 0 | 3257661 | 2456089 | 2500324 | 2452736 | 3005627 | 2545486 | 2544649 | 2159802 | 2496940 |
| 262144 | 0 | 0 | 0 | 0 | 2720150 | 2929641 | 2766395 | 2579573 | 2826531 | 2838190 | 2845962 | 2822482 | 2388073 |
| 524288 | 0 | 0 | 0 | 0 | 3103546 | 3554403 | 3246568 | 3328092 | 3186949 | 3173982 | 3275427 | 3297678 | 3014397 |

2.5 Network Benchmark



A measurement of network speed over the loopback interface card is done for 1 node, between 2 processes on the same node and the network speed between two nodes. The parameter space here includes TCP protocol stack, UDP with fixed packet/buffer size of 64KB. Multiple threads are parallelly run achieve concurrency. Throughput is measured in Mb/sec and Latency in ms for Network benchmarks.

| TCP | No. of Threads | 64KB |
|------------|----------------|-------------|
| Throughput | 1 | 128.2922306 |
| | 2 | 119.6569704 |
| | 4 | 89.6503151 |
| | 8 | 57.09115108 |
| Latency | 1 | 3.990888596 |
| | 2 | 4.278898239 |
| | 4 | 5.711078644 |
| | 8 | 8.968114853 |





| UDP | No. of Threads | 64KB |
|------------|----------------|-------------|
| Throughput | 1 | 162.897948 |
| | 2 | 133.682996 |
| | 4 | 115.549295 |
| | 8 | 79.05041773 |
| Latency | 1 | 3.143072128 |
| | 2 | 3.829956055 |
| | 4 | 4.431009293 |
| | 8 | 6.47687912 |





Multi Nodes (2 nodes used for testing):


Throughput is measured in Mb/sec and Latency in ms for Network benchmarks.


| TCP (2 nodes) | No. of Threads | 64KB |
|---------------|----------------|-------------|
| Throughput | 1 | 98.19761525 |
| | 2 | 83.68993172 |
| | 4 | 74.25600443 |
| | 8 | 50.52308312 |

| | | |
|----------------|---|-------------|
| Latency | 1 | 5.213975906 |
| | 2 | 6.11782074 |
| | 4 | 6.895065308 |
| | 8 | 10.1339817 |



| UDP (2 nodes) | No. of Threads | 64KB |
|---------------|----------------|-------------|
| Throughput | 1 | 80.21079625 |
| | 2 | 109.2866996 |
| | 4 | 94.79908392 |
| | 8 | 68.13947354 |
| Latency | 1 | 6.383180618 |
| | 2 | 4.684925079 |
| | 4 | 5.400896072 |
| | 8 | 7.513999939 |





Comparing to theoretical memory performance, the UDP is different because there is size limitation for packet, no more than 65536 bytes per packet.

For loopback, the entire process is usually CPU bound so its performance is therefore directly linked to CPU speed plus stack efficiency. The traffic never hits the physical network. However, for two nodes, we have to count the physical network in.

We observe that TCP throughput is lesser than UDPs. It is because that UDP is faster than TCP, and the reason is because it's nonexistent acknowledge packet (ACK)(connectionless protocol), instead of TCP that acknowledges a set of packets and has to create a handshake first.

And for two nodes, both of our throughput are lesser than the theoretical network speed. It might be caused by the unknown configuration of Chameleon's Hubs/routers and the latency between two machines(nodes).

In the IPerf results, it has given bandwidth and packets which are transferred across two sides. However, for UDP, it is not feasible to send more than 65536 bytes at the same time due to the limitation of UDP. On the other hand, TCP can send more bytes. So we have to consider this when we look into the results.

There is difference between the theoretical latency. Our conclusion is that even we already keep only one process(our code for testing) running but for CentOS itself there are still some

background processes running. This will be one of the reasons that affect our results of both IPerf and our code.

IPerf Benchmark results are as below:

Results are stored in Network/IPERF_loopback.png and Network/IPERF_2nodes.txt

LOOPBACK:

```
[cc@pal-liu-pereira ~]$ iperf -c 127.0.0.1 -t 20 -p 9999 -w 40k
-----
Client connecting to 127.0.0.1, TCP port 9999
TCP window size: 80.0 KByte (WARNING: requested 40.0 KByte)
-----
[ 3] local 127.0.0.1 port 33268 connected with 127.0.0.1 port 9999
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-20.0 sec  64.6 MBytes 27.1 Mbits/sec
[cc@pal-liu-pereira ~]$ iperf -c 127.0.0.1 -u -p 9999
-----
Client connecting to 127.0.0.1, UDP port 9999
Sending 1470 byte datagrams, IPG target: 11215.21 us (kalman adjust)
UDP buffer size: 208 KByte (default)
-----
[ 3] local 127.0.0.1 port 33068 connected with 127.0.0.1 port 9999
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-10.0 sec  1.25 MBytes 1.05 Mbits/sec
[ 3] Sent 893 datagrams
[ 3] Server Report:
[ 3] 0.0-10.0 sec  1.25 MBytes 1.05 Mbits/sec  0.000 ms  0/ 893 (0%)
[cc@pal-liu-pereira ~]$
```

```
[cc@pal-liu-pereira ~]$ iperf -s -p 9999
-----
Server listening on TCP port 9999
TCP window size: 85.3 KByte (default)
-----
[ 4] local 127.0.0.1 port 9999 connected with 127.0.0.1 port 33268
[ ID] Interval      Transfer    Bandwidth
[ 4] 0.0-20.0 sec  64.6 MBytes 27.1 Mbits/sec
^C[cc@pal-liu-pereira ~]$ iperf -s -u -p 9999
-----
Server listening on UDP port 9999
Receiving 1470 byte datagrams
UDP buffer size: 208 KByte (default)
-----
[ 3] local 127.0.0.1 port 9999 connected with 127.0.0.1 port 33068
[ ID] Interval      Transfer    Bandwidth      Jitter    Lost/Total Datagrams
[ 3] 0.0-10.0 sec  1.25 MBytes 1.05 Mbits/sec  0.000 ms  0/ 893 (0%)
```

cc

0 > cc@pal-liu-pereira:~

pal-liu-pereira

2 nodes:

TCP:

Server:

```
[cc@pa1-lp-t ~]$ iperf -s -p 9999
```

Server listening on TCP port 9999

TCP window size: 85.3 KByte (default)

[4] local 192.168.0.186 port 9999 connected with 192.168.0.178 port 40738

| [ID] | Interval | Transfer | Bandwidth |
|-------|----------|----------|-----------|
|-------|----------|----------|-----------|

| | | | |
|------|--------------|-------------|----------------|
| [4] | 0.0-20.0 sec | 2.33 GBytes | 1.00 Gbits/sec |
|------|--------------|-------------|----------------|

Client:

```
[cc@pa1-liu-pereira ~]$ iperf -c 192.168.0.186 -t 20 -p 9999 -w 40k
```

Client connecting to 192.168.0.186, TCP port 9999

TCP window size: 80.0 KByte (WARNING: requested 40.0 KByte)

[3] local 192.168.0.178 port 40738 connected with 192.168.0.186 port 9999

| [ID] | Interval | Transfer | Bandwidth |
|-------|----------|----------|-----------|
|-------|----------|----------|-----------|

| | | | |
|------|--------------|-------------|----------------|
| [3] | 0.0-20.0 sec | 2.33 GBytes | 1.00 Gbits/sec |
|------|--------------|-------------|----------------|

UDP:

Server:

```
[cc@pa1-lp-t ~]$ iperf -s -u -p 9999
```

Server listening on UDP port 9999

Receiving 1470 byte datagrams

UDP buffer size: 208 KByte (default)

[3] local 192.168.0.186 port 9999 connected with 192.168.0.178 port 58691
[ID] Interval Transfer Bandwidth Jitter Lost/Total Datagrams
[3] 0.0-10.0 sec 1.25 MBytes 1.05 Mbits/sec 0.019 ms 0/ 893 (0%)

Client:

```
[cc@pa1-liu-pereira ~]$ iperf -c 192.168.0.186 -u -p 9999
```

Client connecting to 192.168.0.186, UDP port 9999

Sending 1470 byte datagrams, IPG target: 11215.21 us (kalman adjust)

UDP buffer size: 208 KByte (default)

[3] local 192.168.0.178 port 58691 connected with 192.168.0.186 port 9999
[ID] Interval Transfer Bandwidth
[3] 0.0-10.0 sec 1.25 MBytes 1.05 Mbits/sec
[3] Sent 893 datagrams
[3] Server Report:
[3] 0.0-10.0 sec 1.25 MBytes 1.05 Mbits/sec 0.019 ms 0/ 893 (0%)
